



Video 1.2

Chris Murphy

Review

- The **Internet** is a physical network of devices
- The **World Wide Web** is an application that utilizes the Internet to allow for accessing data
- Resources on the Web have unique **URLs** that include the protocol, host name, and file/resource name

#<<<>>>
#copyright

Your continued donations keep Wikipedia running!

Lynx (web browser)

From Wikipedia, the free encyclopedia

Jump to: [navigation](#), [search](#)

CAPTION: Lynx

Wikipedia Main Page displayed in Lynx

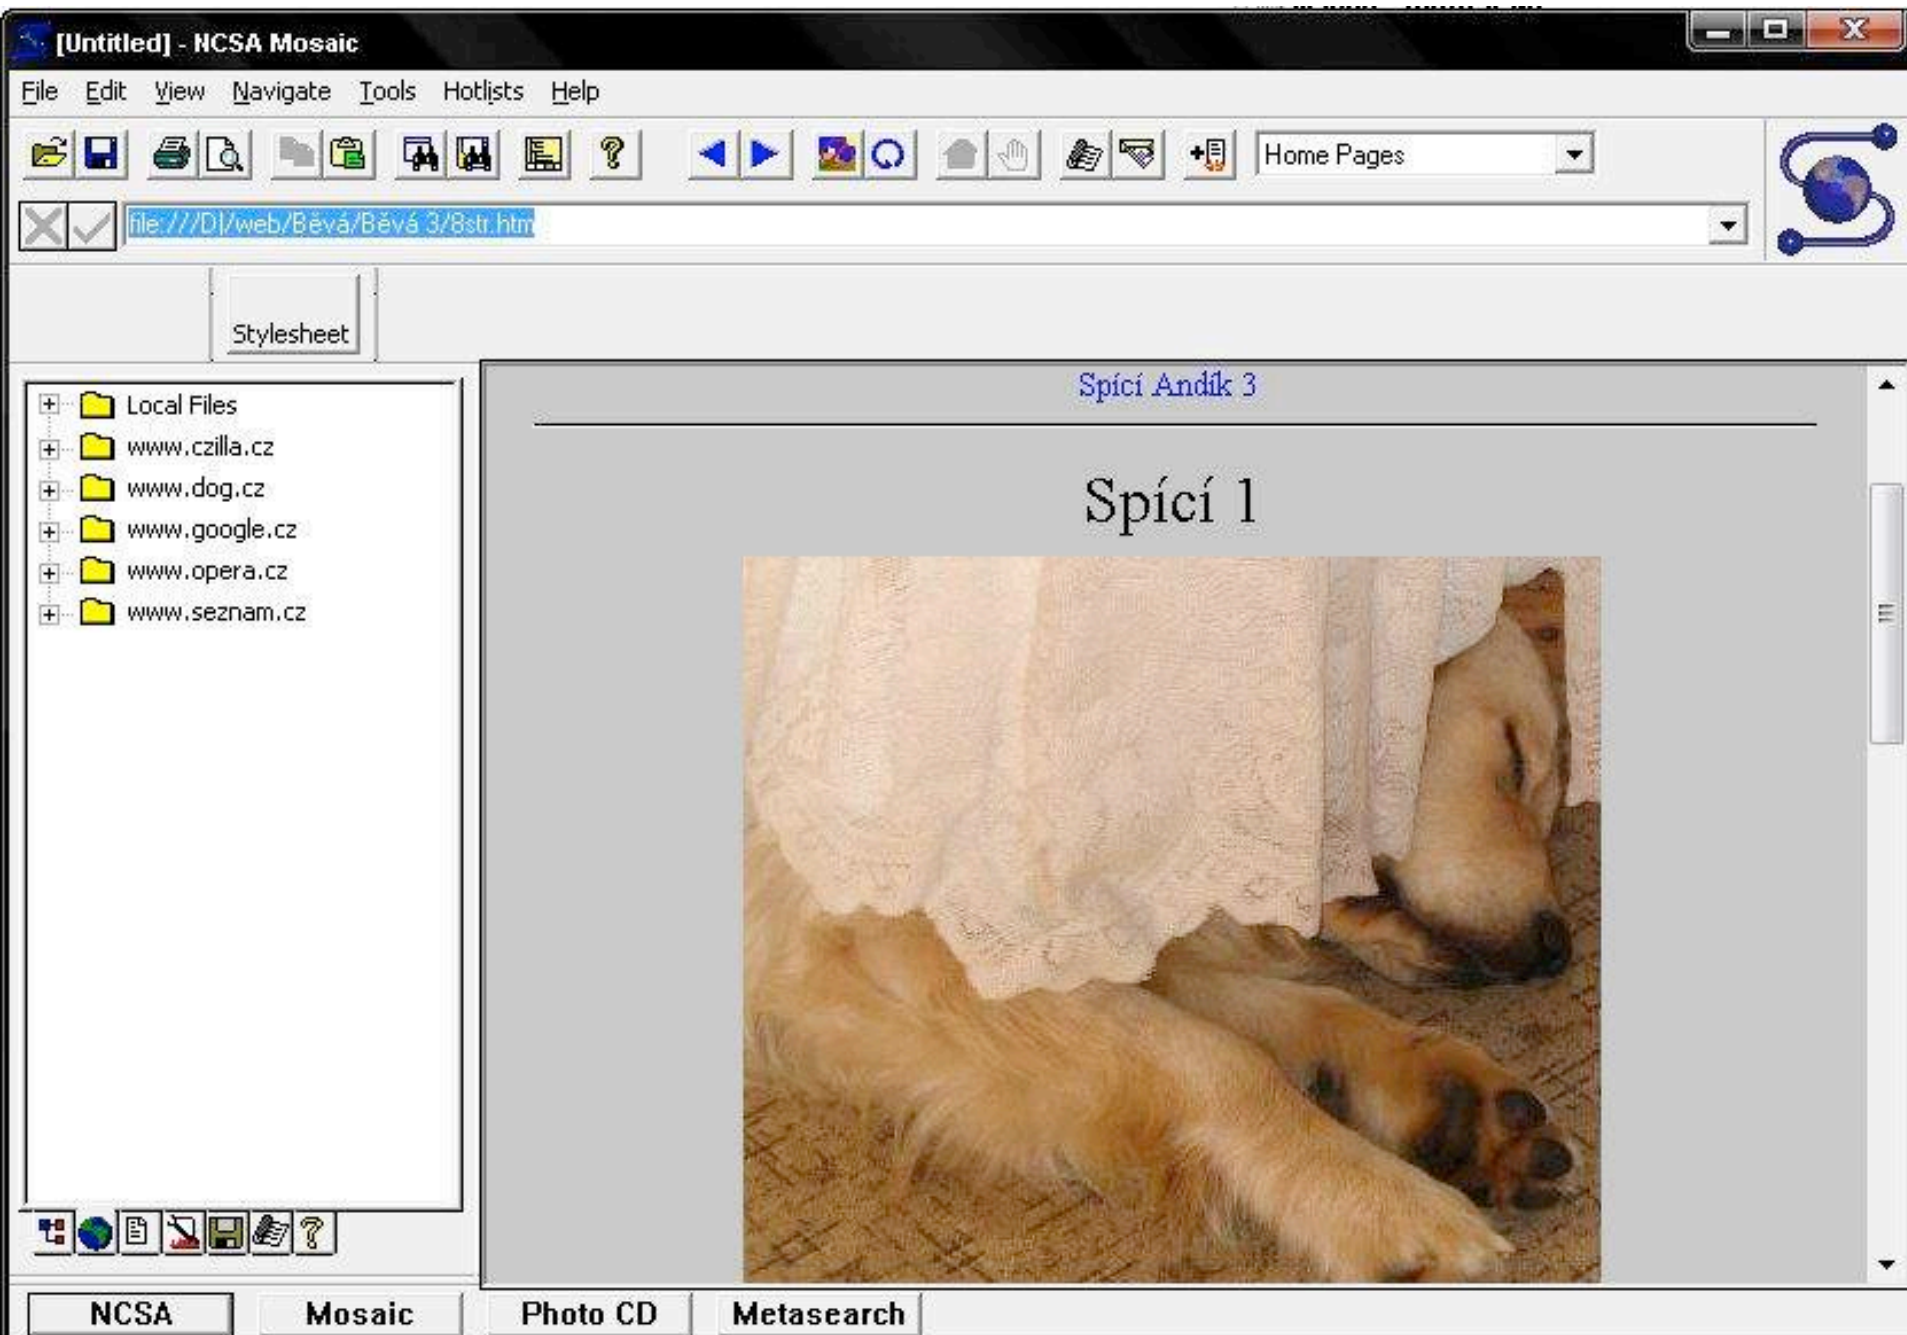
Wikipedia Main Page displayed in Lynx

Maintainer: Thomas Dickey
Stable release: 2.8.5 (February 4, 2004) [\[\[+/-\]\]](#)
Preview release: 2.8.6 (?) [\[\[+/-\]\]](#)
OS: Cross-platform
Use: web browser
License: GPL
Website: lynx.isc.org

Lynx is a text-only **Web browser** and **Internet Gopher** client for use on cursor-addressable, character cell **terminals**.

Browsing in Lynx consists of highlighting the chosen link using cursor keys, or having all links on a page numbered and entering the chosen link's number. Current versions support **SSL** and many **HTML** features. Tables are linearized (scrunched together one cell after another without tabular structure), while frames are identified by name and can be explored as if they were separate pages.

Lynx is a product of the Distributed Computing Group within Academic Computing Services of the **University of Kansas**, and was initially developed in 1992 by a team of students at the university (Lou Montulli, Michael Grobe and Charles Rezac) as a hypertext browser used solely to distribute campus information as part of a **Campus-Wide Information Server**. In 1993 Montulli added an Internet interface and released a new version (2.0) of the browser [\[1\]](#) [\[2\]](#) [\[3\]](#).



Popular Browsers



What is a Web Browser?

- **Browser:** software that is used to access and display Web content, and to navigate across the Web
- **Main Components of the Browser**
 - Rendering Engine (HTML/CSS) – responsible for static content presentation, formatting, and layout
 - JavaScript Engine (JavaScript) – responsible for creating and modifying dynamic content and appearance

How Does a Web Browser Work?

- Browser and the World Wide Web utilize **Hypertext Transfer Protocol (HTTP)** to transfer documents

How Does a Web Browser Work?

- Browser and the World Wide Web utilize **Hypertext Transfer Protocol (HTTP)** to transfer documents



Client

How Does a Web Browser Work?

- Browser and the World Wide Web utilize **Hypertext Transfer Protocol (HTTP)** to transfer documents



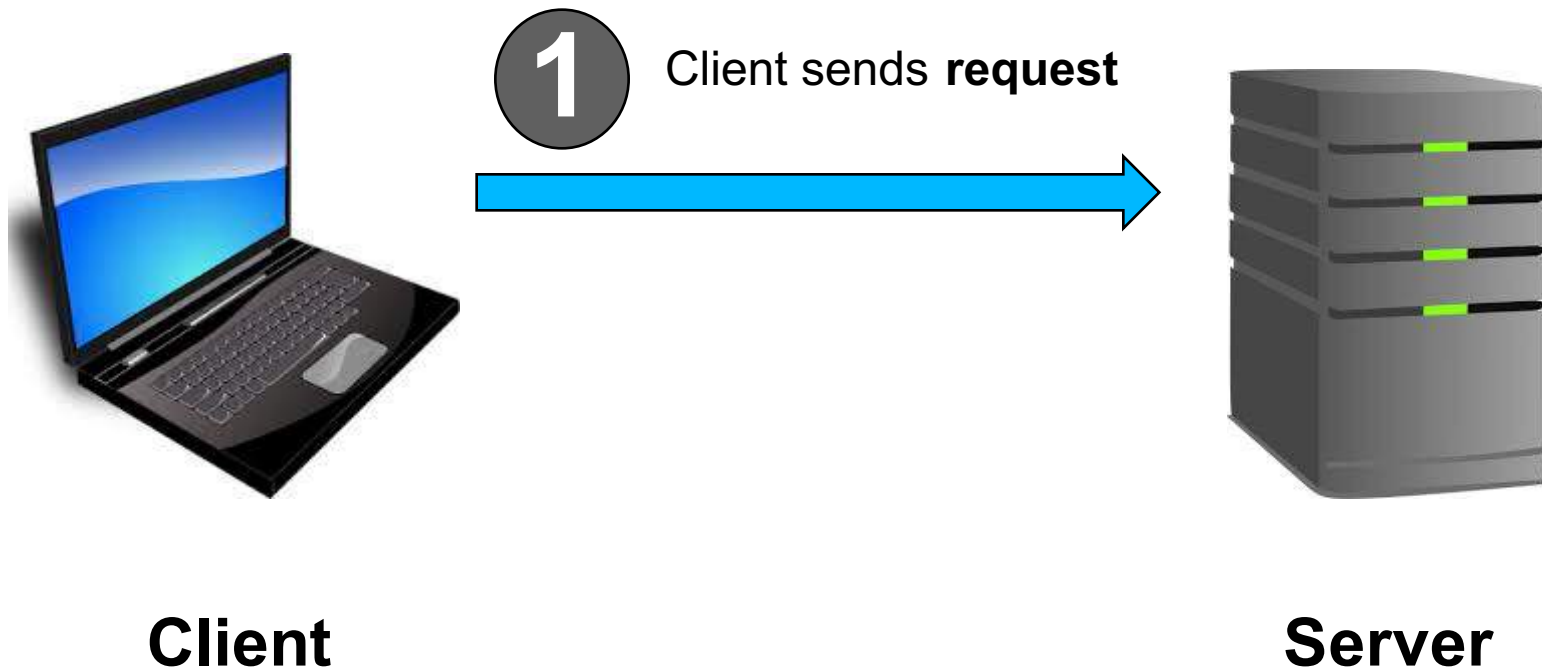
Client



Server

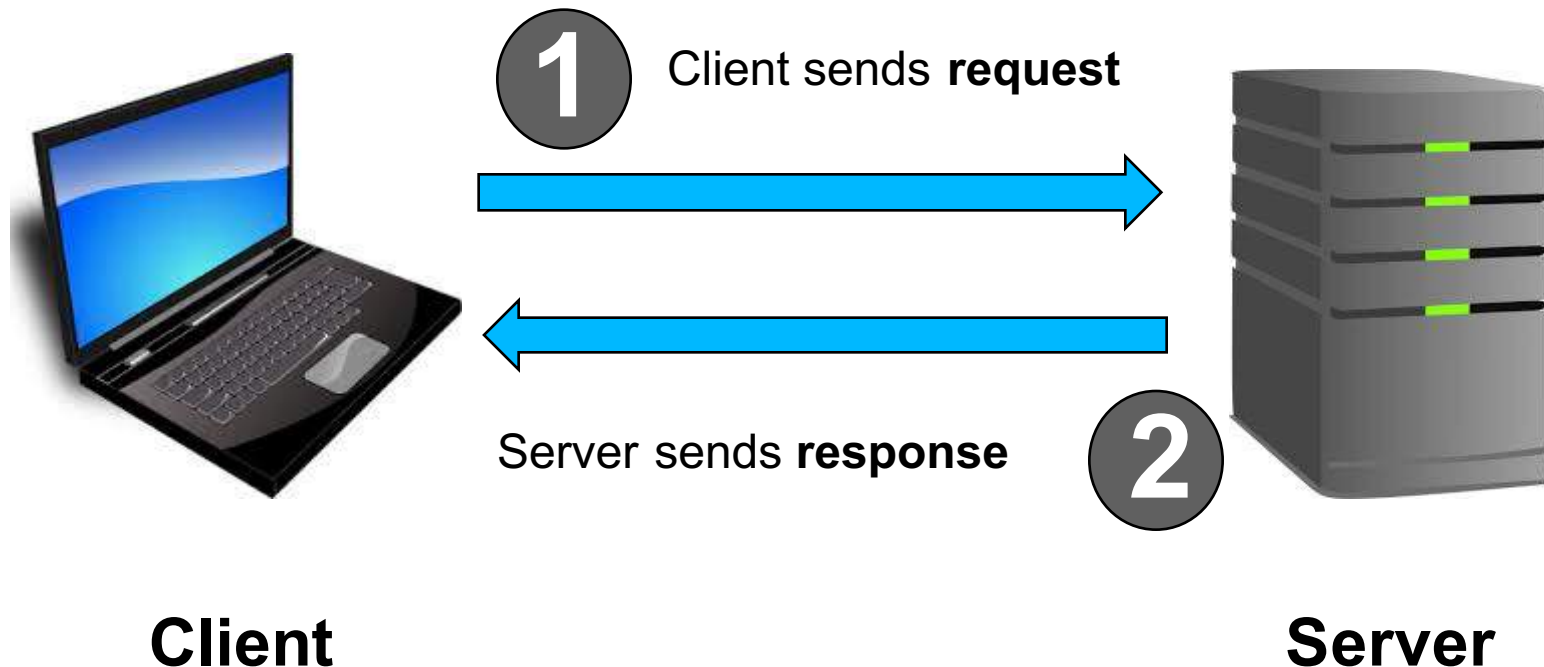
How Does a Web Browser Work?

- Browser and the World Wide Web utilize **Hypertext Transfer Protocol (HTTP)** to transfer documents



How Does a Web Browser Work?

- Browser and the World Wide Web utilize **Hypertext Transfer Protocol (HTTP)** to transfer documents



HTTP Overview

- HTTP is a plain-text, human-readable protocol used for exchanging data on the Web
- Initially developed by Tim Berners-Lee at CERN in 1989
- Based on client-server model:
 - Client sends **request** for resource, possibly including information about the client
 - Server sends **response**, including header (status information) and requested resource

```
josh@blackbox:~$ telnet en.wikipedia.org 80
Trying 208.80.152.2...
Connected to rr.pmtpa.wikimedia.org.
Escape character is '^]'.
GET /wiki/Main_Page http/1.1
Host: en.wikipedia.org
```

Request

```
HTTP/1.0 200 OK
Date: Thu, 03 Jul 2008 11:12:06 GMT
Server: Apache
X-Powered-By: PHP/5.2.5
Cache-Control: private, s-maxage=0, max-age=0, must-revalidate
Content-Language: en
Vary: Accept-Encoding, Cookie
X-Vary-Options: Accept-Encoding;list-contains=gzip, Cookie;string-contains=enwikiToken;string-contains=enwikiLoggedOut;string-contains=enwiki_session;
string-contains=centralauth_Token;string-contains=centralauth_Session;string-contains=centralauth_LoggedOut
Last-Modified: Thu, 03 Jul 2008 10:44:34 GMT
Content-Length: 54218
Content-Type: text/html; charset=utf-8
X-Cache: HIT from sq39.wikimedia.org
X-Cache-Lookup: HIT from sq39.wikimedia.org:3128
Age: 3
X-Cache: HIT from sq38.wikimedia.org
X-Cache-Lookup: HIT from sq38.wikimedia.org:80
Via: 1.0 sq39.wikimedia.org:3128 (squid/2.6.STABLE18), 1.0 sq38.wikimedia.org:80 (squid/2.6.STABLE18)
Connection: close
```

Response headers

```
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Transitional//EN" "http://www.w3.org/TR/xhtml1/DTD/xhtml1-transitional.dtd">
<html xmlns="http://www.w3.org/1999/xhtml" xml:lang="en" lang="en" dir="ltr">
  <head>
    <meta http-equiv="Content-Type" content="text/html; charset=utf-8" />
    <meta name="keywords" content="Main Page,1778,1844,1863,1938,1980 Summer Olympics,2008,2008 Guizhou riot,2008 Jerusal
...
... This content has been removed to save space
...
"Non-profit organization">nonprofit</a> <a href="http://en.wikipedia.org/wiki/Charitable_organization" title="Charitable organization">charity</a>.<b
r /></li>
    <li id="privacy"><a href="http://wikimediafoundation.org/wiki/Privacy_policy" title="wikimedia:Privacy policy">Privac
y policy</a></li>
    <li id="about"><a href="/wiki/Wikipedia:About" title="Wikipedia:About">About Wikipedia</a></li>
    <li id="disclaimer"><a href="/wiki/Wikipedia:General_disclaimer" title="Wikipedia:General disclaimer">Disclaimers</a>
</li>
  </ul>
</div>
</div>
<script type="text/javascript">if (window.runOnloadHook) runOnloadHook();</script>
<!-- Served by srv93 in 0.050 secs. --></body></html>
```

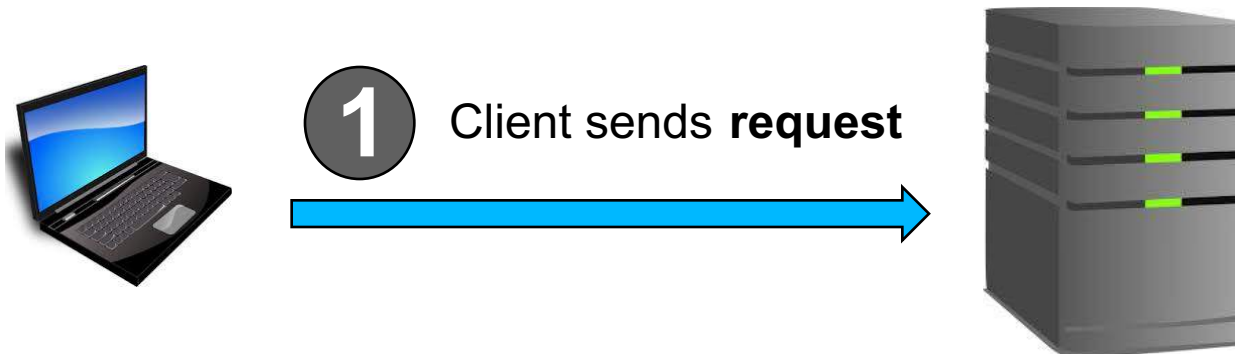
Response body

Connection closed by foreign host.

```
josh@blackbox:~$
```

Anatomy of an HTTP Request

- First line of request will always be a **verb** followed by an **argument**
 - **GET** – retrieve resource
 - **HEAD** – retrieve only headers (information about the resource)
 - **POST** – create resource (usually used in form submission context)
- Next comes the protocol (usually HTTP/1.1)
- Optionally include other information about the request and/or the client



HTTP Request Example

```
GET /examples/index.html HTTP/1.1
Host: www.edx.org
User-Agent: Mozilla/4.0
Accept-Language: en-us
Content-Length: 9

a=12&b=34
```

HTTP Request Example

```
GET /examples/index.html HTTP/1.1
Host: www.edx.org
User-Agent: Mozilla/4.0
Accept-Language: en-us
Content-Length: 9

a=12&b=34
```

→ Request Line

HTTP Request Example

Request Verb

```
GET /examples/index.html HTTP/1.1  
Host: www.edx.org  
User-Agent: Mozilla/4.0  
Accept-Language: en-us  
Content-Length: 9  
  
a=12&b=34
```

Request Line

HTTP Request Example

Request Verb

URI

```
GET /examples/index.html HTTP/1.1
Host: www.edx.org
User-Agent: Mozilla/4.0
Accept-Language: en-us
Content-Length: 9

a=12&b=34
```

Request Line

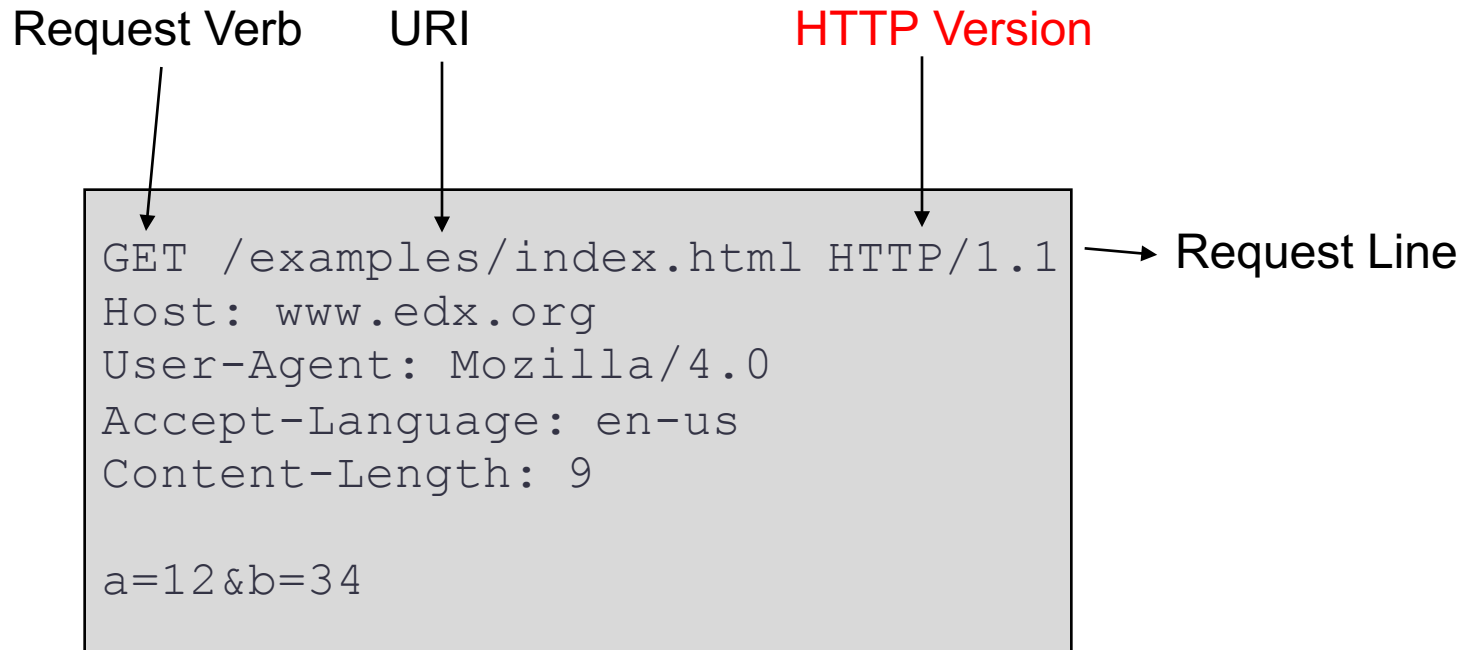
HTTP Request Example

Request Verb URI HTTP Version

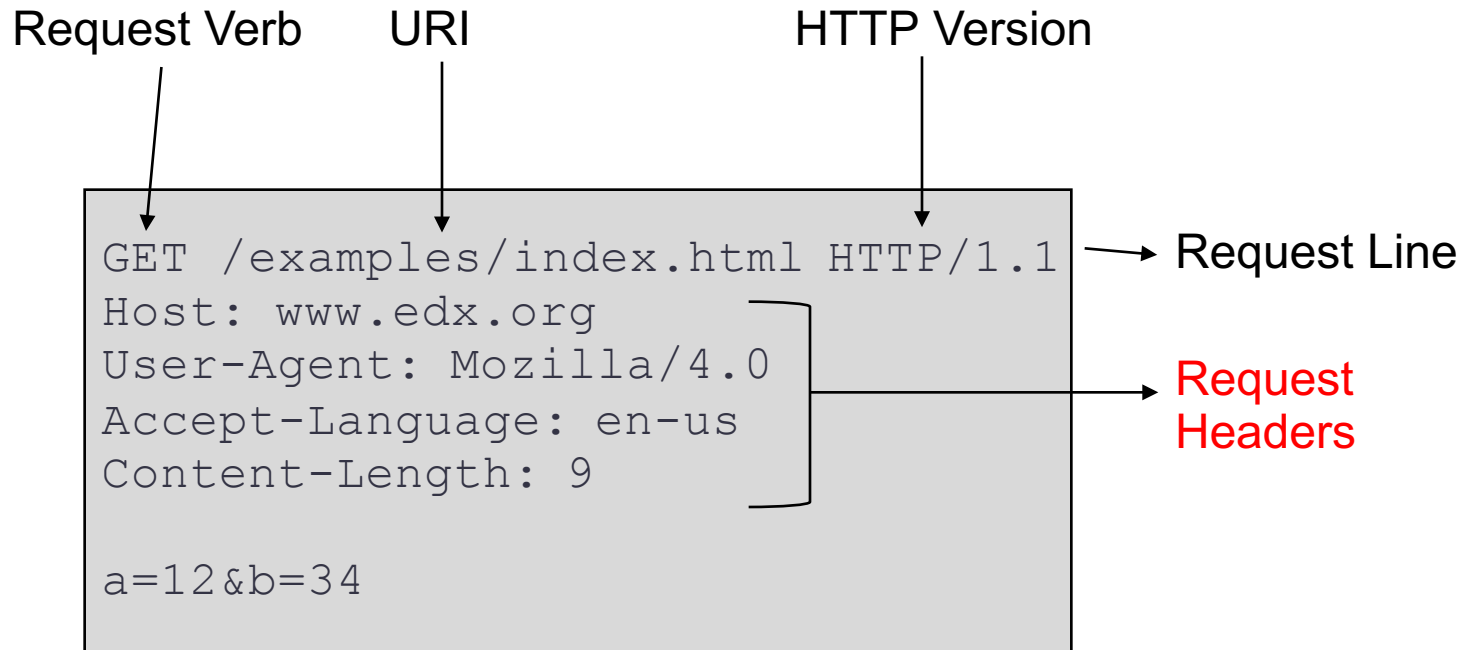
```
GET /examples/index.html HTTP/1.1
Host: www.edx.org
User-Agent: Mozilla/4.0
Accept-Language: en-us
Content-Length: 9

a=12&b=34
```

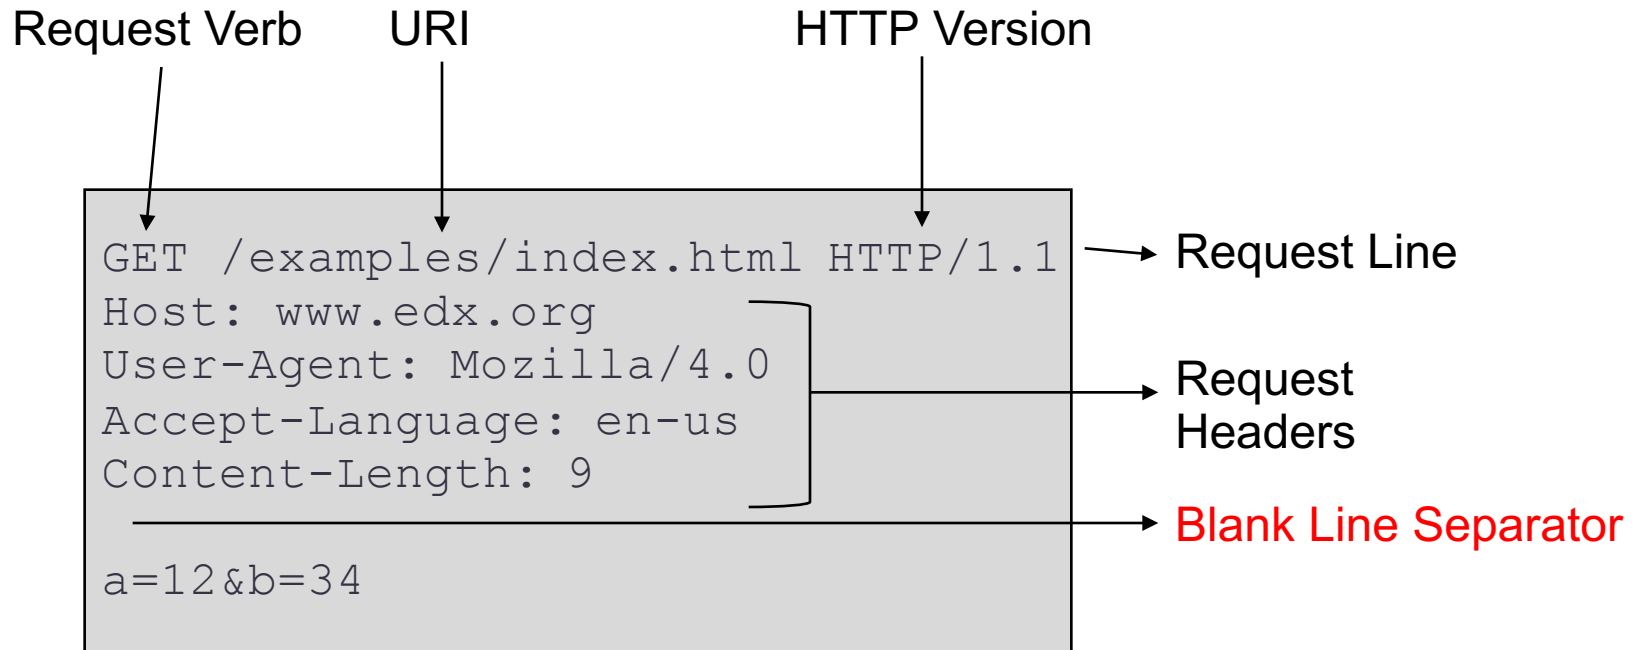
Request Line

The diagram shows an HTTP request structure. Above a gray box containing the request, three labels are positioned: 'Request Verb' above 'GET', 'URI' above '/examples/index.html', and 'HTTP Version' above 'HTTP/1.1'. Arrows point from each label to its corresponding part in the request line. To the right of the box, the label 'Request Line' has an arrow pointing to the first line of the request, 'GET /examples/index.html HTTP/1.1'. The box also contains the Host, User-Agent, Accept-Language, Content-Length headers, and a body with query parameters.

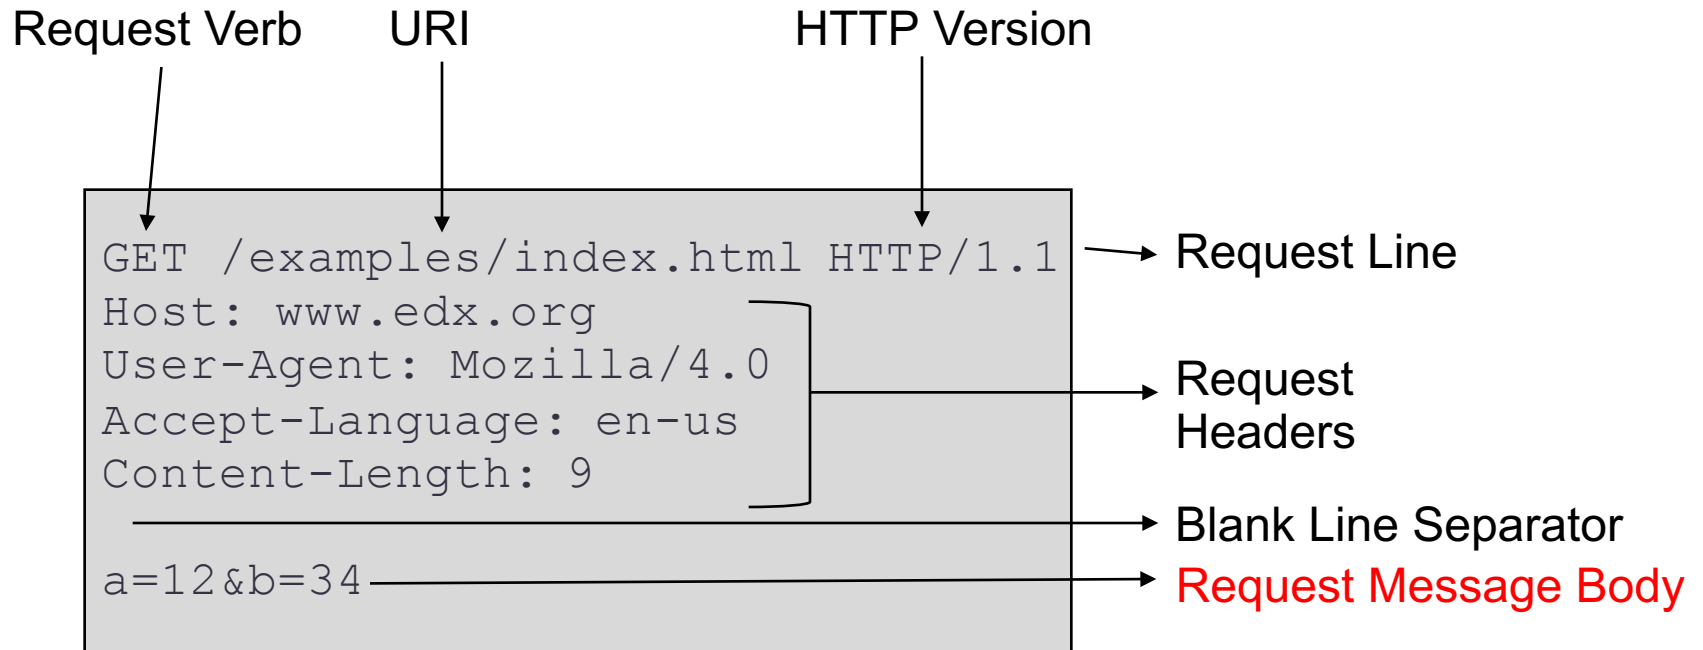
HTTP Request Example



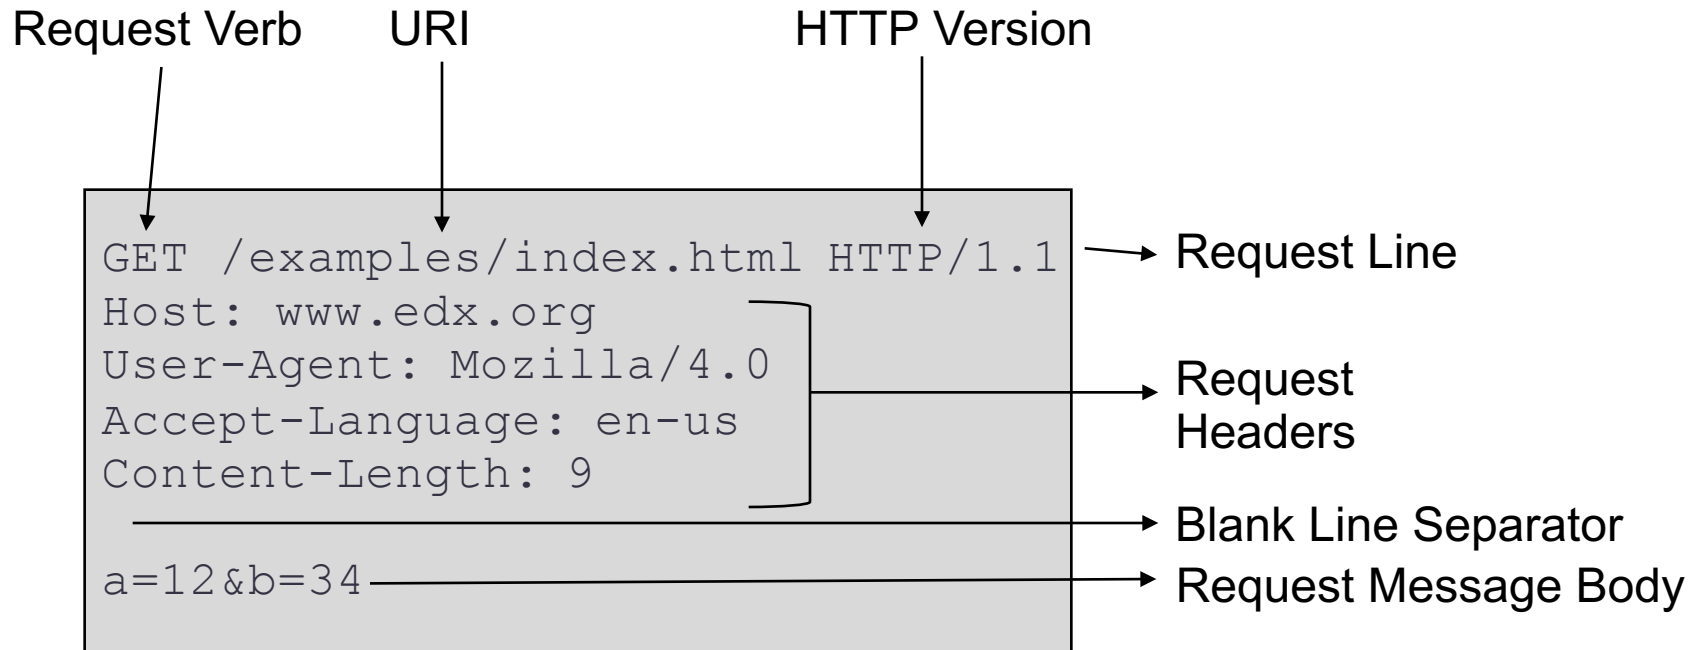
HTTP Request Example



HTTP Request Example



HTTP Request Example



Anatomy of an HTTP Response

- First line is always protocol and **status code**
 - 1XX – information only
 - 2XX – success
 - 3XX – client redirect
 - 4XX – client error
 - 5XX – server error



Server sends **response**

2



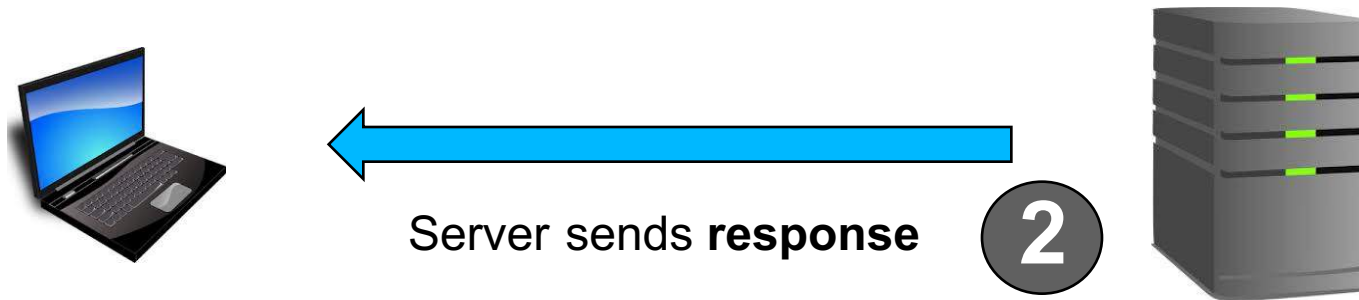
Most Common Status Codes

- **200 OK** – request succeeded, resulting resource (as stated in request) will be included in message body
- **404 Not Found** – requested resource does not exist
- **500 Server Error** – Error on the server side in processing request



Anatomy of an HTTP Response

- Following protocol and status code will be other **header information** regarding the response and/or the server
- Then a blank line
- Then the response body, i.e. the resource that was requested



HTTP Response Example

```
HTTP/1.1 200 OK
Date: Fri, 06 Apr xxxx 09:30:00 GMT
Server: Apache/1.4
Last-Modified: Wed, 04 Apr xxxx
Connection: close
Content-Type: text/html
Content-Length: 228

<!DOCTPYE html><html><head>...
```

HTTP Response Example

```
HTTP/1.1 200 OK
Date: Fri, 06 Apr xxxx 09:30:00 GMT
Server: Apache/1.4
Last-Modified: Wed, 04 Apr xxxx
Connection: close
Content-Type: text/html
Content-Length: 228

<!DOCTPYE html><html><head>...
```

→ Response Line

HTTP Response Example

HTTP Version

```
HTTP/1.1 200 OK
Date: Fri, 06 Apr xxxx 09:30:00 GMT
Server: Apache/1.4
Last-Modified: Wed, 04 Apr xxxx
Connection: close
Content-Type: text/html
Content-Length: 228
```

Response Line

```
<!DOCTYPE html><html><head>...
```

HTTP Response Example

HTTP Version

Status Code

```
HTTP/1.1 200 OK
Date: Fri, 06 Apr xxxx 09:30:00 GMT
Server: Apache/1.4
Last-Modified: Wed, 04 Apr xxxx
Connection: close
Content-Type: text/html
Content-Length: 228
```

```
<!DOCTPYE html><html><head>...
```

Response Line

HTTP Response Example

HTTP Version

Status Code

```
HTTP/1.1 200 OK
Date: Fri, 06 Apr xxxx 09:30:00 GMT
Server: Apache/1.4
Last-Modified: Wed, 04 Apr xxxx
Connection: close
Content-Type: text/html
Content-Length: 228
```

```
<!DOCTYPE html><html><head>...
```

Response Line

Response
Headers

HTTP Response Example

HTTP Version

Status Code

```
HTTP/1.1 200 OK
Date: Fri, 06 Apr xxxx 09:30:00 GMT
Server: Apache/1.4
Last-Modified: Wed, 04 Apr xxxx
Connection: close
Content-Type: text/html
Content-Length: 228
```

```
<!DOCTYPE html><html><head>...
```

Response Line

Response
Headers

Blank Line Separator

HTTP Response Example

HTTP Version

Status Code

```
HTTP/1.1 200 OK
Date: Fri, 06 Apr xxxx 09:30:00 GMT
Server: Apache/1.4
Last-Modified: Wed, 04 Apr xxxx
Connection: close
Content-Type: text/html
Content-Length: 228
```

```
<!DOCTYPE html><html><head>...
```

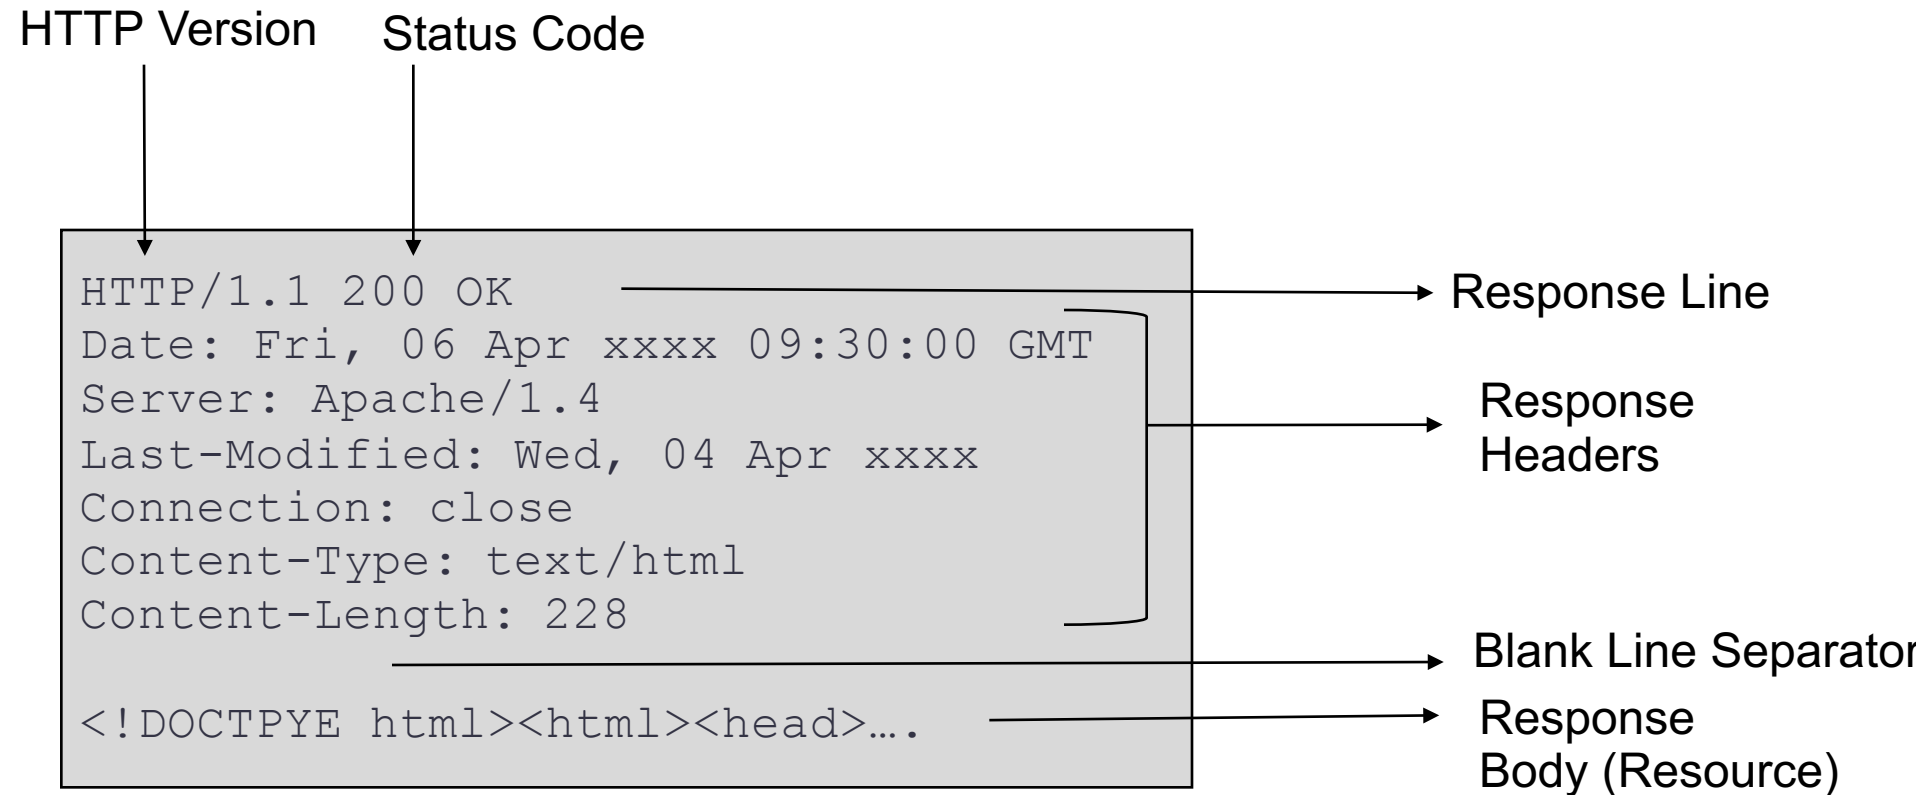
Response Line

Response
Headers

Blank Line Separator

Response
Body (Resource)

HTTP Response Example



Summary

- Web browsers are used to access data on the Web
- Browsers communicate with web servers using HTTP
- HTTP is based on a client-server model:
 - Client sends **request** for resource, possibly including information about the client
 - Server sends **response**, including header (status information) and requested resource